


A 256-token budget contains less than two dialogue turns

Assumption-eligible unique source contexts; episode-clustered 95% confidence interval



95% CI: 1.78–1.85 turns

Mean raw-history use: 250.5 of 256 tokens